

Repairing Document Knowledge Graphs with Source Text and Field Checks

Hua Peng^{1†}, Zifu Tao^{2†}, Yu Guo¹, Qiang Zhou², Tian Zhou^{3*},
Bo Han³

¹China Mobile System Integration Co. Ltd., China.

²Tongji University, China.

³Xi'an Jiaotong University, China.

*Corresponding author(s). E-mail(s): tianzhou@xjtu.edu.cn;

Contributing authors: penghua@cmict.chinamobile.com;

2253323@tongji.edu.cn; guoyu1@cmict.chinamobile.com;

zhouqiang@cmict.chinamobile.com; bohan@xjtu.edu.cn;

[†]These authors contributed equally to this work.

Abstract

Knowledge graphs built from documents can have missing fields and wrong values. We repair these graphs using the source text and field rules. The method cleans the input, asks a language model for a graph, and checks each triple. The checks cover the document ID, field name, source text, and number of values per field. We use shared prompts and score each response before and after filtering. On 225 graphs with 450 added defects, Diagnosis + Gate repairs 98.00% of defects. On 225 graphs built by a model, it raises triple F1 from 87.99% to 92.48%. Direct field copying reaches 100% on these records. On a new test of 60 receipts, a source-line index raises F1 over Simple Pipeline by 4.76 points (95% CI: 2.50 to 7.26). On 56 CUAD contracts, indexed repair reaches 44.84% F1, compared with 46.76% for the input. Its F1 difference from indexed fresh extraction is 0.72 points (95% CI: **-4.80** to 6.48). Two reviewers check 200 reference differences and 200 system edits. They find input errors in 95 of the 200 differences. After a third reviewer resolves disagreements, 134 of 199 edits from five settings are accepted. One edit remains uncertain. The tests show how source text, field checks, and input graphs affect repair. Code, model outputs, and review rules are available for further study.

Keywords: Knowledge Graph Repair, Source Text, Field Checks, Document Extraction

1 Introduction

A knowledge graph (KG) stores facts as triples. Each triple has a subject, a relation, and a value. For example, a receipt graph can store a receipt ID, the relation “total”, and the value “120.00”. Graphs built from text often have missing values, wrong values, repeated triples, or invalid relations. These errors affect search and later use of the data (Ji et al. 2022; Paulheim 2017). The source text gives us a way to check and repair the graph.

Field rules and language models play two roles. Rules check the graph’s form. A model reads the text and proposes field values. KG quality tools and SHACL provide rules for graph checks (Zaveri et al. 2016; Kontokostas et al. 2014; Knublauch and Kontokostas 2017). Lin et al. add rules and graph context to LLM repair prompts (Lin et al. 2025). We study three questions: when does repair improve a graph, which steps produce the gain, and when does an existing graph help more than building a new one?

We study graphs that store the fields of one document. Each graph has a fixed document ID and a list of allowed relations. Each relation has at most one value. Our method first cleans the input graph. A language model then reads the source text and returns a complete graph. It can also read an error report. A filter checks the document ID, allowed relations, text matches, and field counts. It saves a reason for each rejected triple.

We compare three prompt settings: basic input, input with an error report, and input with SHACL checks. They share the same instructions and output limit. We score each response before and after filtering. This shows the change from extra prompt text and the change from the filter. We also test a selector that chooses edits one at a time and measure the cost of graph checks.

The tests use three Chinese record collections, SROIE receipts, and CUAD contracts. Chinese records list field names and values. Receipt text and human field labels are stored apart. We use 20 new receipts to set the source-line index, then fix the design and test it on 60 further receipts. The contract test uses 56 documents. Two reviewers also check 200 reference differences and 200 system edits. These tests measure field recovery, the loss of correct fields, and human acceptance of edits.

The paper makes three contributions:

- A repair method that joins source-based generation with field checks and records why each triple is rejected.
- Tests on added defects, model outputs, receipts, and contracts, with scores for field recovery and retention and a human review of errors and edits.
- Shared-prompt and shared-response tests of error reports, source-line hints, SHACL context, and filtering. We also report full edit-selector runs and release the outputs and logs.

Section 2 reviews earlier work. Sections 3 and 4 give the method and its implementation. Section 5 gives the tests and results. Section 6 concludes.

2 Related Work

2.1 Building Knowledge Graphs

Building a KG involves extracting facts, linking entities, mapping relations, and joining the results (Ji et al. 2022). DBpedia (Auer et al. 2007), YAGO (Suchanek et al. 2007), and Freebase (Bollacker et al. 2008) build knowledge bases from data with clear or partly clear structure. Knowledge fusion uses source evidence to score facts (Dong et al. 2014). Our study repairs graphs after extraction.

Embedding models use vectors to predict links (Bordes et al. 2013; Wang et al. 2017). Graph neural networks learn from nearby nodes and relations (Schlichtkrull et al. 2018). Reviews describe the use of LLMs to build, complete, and combine KGs (Pan et al. 2023, 2024; Bian 2025). Prompt methods include chain-of-thought reasoning (Wei et al. 2022) and shared prompt patterns (White et al. 2023). Retrieval-augmented generation adds source evidence to model inputs (Lewis et al. 2020). Neo4j Labs’ LLM Knowledge Graph Builder turns documents into graphs with user-set schemas (Neo4j Labs 2024). We compare with its core extractor in Section A.7.

2.2 Checking Graph Quality

KG quality covers correct facts, missing facts, rule agreement, and current information (Zaveri et al. 2016; Paulheim 2017). Statistical methods find unusual values (Wienand and Paulheim 2014; Paulheim and Bizer 2014). Embedding models score how likely a triple is (Nickel et al. 2016). Rule languages check the graph’s schema (Kontokostas et al. 2014). Each method checks a different part of graph quality.


Lin et al. add rule errors to graphs and compare LLM repair prompts (Lin et al. 2025). They use SHACL rules and nearby graph context. Human review helps check cases with unclear answers (Acosta et al. 2013). Other reviews cover quality checks during graph building and repair (Wang et al. 2021; Paulheim 2017). Our error report lists duplicate triples, schema errors, and values that are missing from the source text.

2.3 Repairing Graphs and Choosing Actions

KG completion predicts missing links. Methods use paths (Lao et al. 2011), embeddings (Bordes et al. 2013; Yang et al. 2015; Shi and Weninger 2017), and neural networks (Dettmers et al. 2018). MINERVA learns to follow graph paths to answer queries (Das et al. 2018). MetaR predicts links from a few examples (Chen et al. 2019). AMIE learns positive rules (Galárraga et al. 2013). RuDiK learns positive and negative rules (Ortona et al. 2018). Other methods add rules to embedding training (Guo et al. 2016; Ding et al. 2018).

RNNLogic trains a rule generator and an answer predictor together (Qu et al. 2021). NBFNet learns from graph paths to predict links (Zhu et al. 2021). pLogicNet joins logical rules and embeddings (Qu and Tang 2019). These methods combine rules with learned predictions. We combine source text with field checks and compare filtering with a selector that scores trial edits.

Lin et al. compare SHACL-based repair prompts (Lin et al. 2025). ReAct joins model reasoning with actions that gather information (Yao et al. 2023). We use this



Sequential variant: reassess after each accepted edit; stop if no feasible positive-utility edit remains.

Fig. 1 Repair by choosing edits one at a time. Four quality scores and four graph counts feed the learned router. It gives a repair score and weights for three kinds of edits. The selector tries model and rule edits on graph copies. It applies the allowed edit with the highest positive score, then checks the graph again. Section 4.4 gives this test.

pattern in the ReAct-style baseline. Our SHACL-context baseline adapts Lin et al.’s five prompt parts to the same field task and call limit. Shared prompts and scores before and after filtering show the effects of context and graph checks.

3 Task and Method

The input is a list of triples G , source text X , a fixed document ID h_0 , and allowed relations $\mathcal{R}_X$. The list can contain copies of the same triple. Each relation has at most one nonempty value. Examples include document fields and receipt fields. The scorer uses the reference graph G^* to check the output.

The main method has three steps: clean the input, ask a model for a complete graph, and filter the triples. The prompt can include an error report or source-line hints. The filter checks each value against the source text. Reference scores and human review measure output quality.

Figures 1 and 3 show a second path. A learned router gives weights to edits. A selector tries the edits on graph copies and chooses one at a time. Section 4.4 tests this path with saved proposals.

3.1 Input Cleanup and Error Report

Figure 2 groups the checks into local, graph, and source checks. They count copies, check the document ID and relation, and look for the value in the text.

The preprocessor P removes duplicate triples. It reverses an edge if its object is h_0 and its subject is a different node. It then removes triples with an invalid head




Fig. 2 Three groups of graph checks: repeated facts and unlinked nodes (local), invalid relations (graph), and values missing from the text (source). Both repair paths use these groups. Their report follows Eq. (1).

or relation. Here, the head is the subject of a triple. The cleaned graph is $G_s = P(G; h_0, \mathcal{R}_X)$.

For a triple $t = (h, r, o)$, the three checks are

$$H(t) = \mathbf{1}[h = h_0], \quad R(t) = \mathbf{1}[r \in \mathcal{R}_X], \quad S(t; X) = \mathbf{1}[o \neq \emptyset \wedge o \text{ occurs in } X]. \quad (1)$$

H checks the document ID, R checks the relation, and S checks whether the nonempty value occurs in the text. The error report stores duplicate counts, failed-check positions, and relations in G_s :

$$D(G_s, X) = (n_{\text{dup}}, I_H, I_R, I_S, \mathcal{R}(G_s)). \quad (2)$$

For example, $I_S = \{i : S(t_i; X) = 0\}$ lists positions that fail the source check. The report uses the cleaned input. Cleanup removes some errors before the report is built. We test the effect of adding this report to the prompt. Fields are optional, and the main method calls the model once for every document to recover missing values.

3.2 Generating a Graph from Source Text

The language model returns an ordered list of triples:

$$C = L_\theta(X, h_0, \mathcal{R}_X, G_s, D(G_s, X)). \quad (3)$$

The prompt asks for a complete graph. It tells the model to copy values from the text, keep correct fields, and fill missing or wrong values. The JSON parser returns an empty graph for an invalid response. Scoring includes these empty graphs. The base prompt uses the same source, input graph, model settings, and output limit, with the error report removed.

3.3 Checking the Generated Triples

The gate is a filter. It keeps triples that pass Eq. (1) and keeps at most one value per relation. Let $\text{uniq}(C)$ remove copies in output order. The triples that pass the checks are

$$C_X = \{t \in \text{uniq}(C) : H(t)R(t)S(t; X) = 1\}. \quad (4)$$

The allowed output sets are

$$\mathcal{F}(C_X) = \left\{ A \subseteq C_X : \sum_{t \in A} \mathbf{1}[r_t = r] \leq 1, \forall r \in \mathcal{R}_X \right\}. \quad (5)$$

The filter scans the triples in order. For each relation, it keeps the first triple that passes the checks. The output $\hat{G}$ has size

$$|\hat{G}| = \max_{A \in \mathcal{F}(C_X)} |A| = |\{r : \exists t \in C_X, r_t = r\}|. \quad (6)$$

Each relation forms one group with room for one triple. The filter keeps one triple from each group that has a valid candidate. It thus keeps the largest allowed number of candidates. The first valid value in output order wins when several values pass.

3.4 Field Meanings and Source Lines

A receipt can list a brand name, a company name, a subtotal, and a final amount. We give each field a short meaning to help the model choose its value. All receipt repair methods receive these meanings and numbered source lines.

The source-line index adds hints for each field. Fixed patterns find company terms, address terms, postcodes, dates, and total labels. Each matching line is an anchor. The index includes that line, one line before it, and two lines after it. Simple receives the full source and field meanings. Evidence Index also receives the line hints.

The receipt filter uses the same whitespace rule as the parser. $N(s)$ changes each run of spaces or line breaks to one space and trims the ends. The source check is

$$S_{\text{ws}}(t; X) = \mathbf{1}[|N(o)| > 0 \wedge N(o) \text{ occurs in } N(X)]. \quad (7)$$

It keeps word boundaries, case, and punctuation. The receipt follow-up uses this check. Earlier tests use the original exact-text check.

Algorithm 1 Repair with source text and field checks

Require: Input triples G , source X , document ID h_0 , allowed relations $\mathcal{R}_X$

Ensure: Repaired graph $\hat{G}$ and rejection log J

```
1:  $G_s \leftarrow P(G; h_0, \mathcal{R}_X)$ 
2:  $D \leftarrow D(G_s, X)$ 
3:  $C \leftarrow \text{PARSE}(L_\theta(X, h_0, \mathcal{R}_X, G_s, D))$ 
4:  $\hat{G} \leftarrow \emptyset; J \leftarrow \emptyset; B \leftarrow \emptyset$ 
5: for each candidate  $t = (h, r, o)$  in generation order do
6:   if  $t$  appeared earlier in  $C$  then
7:     append  $(t, \text{duplicate})$  to  $J$ 
8:   else if  $H(t)R(t)S(t; X) = 1$  and  $r \notin B$  then
9:      $\hat{G} \leftarrow \hat{G} \cup \{t\}; B \leftarrow B \cup \{r\}$ 
10:  else
11:    append  $(t, \text{failed check})$  to  $J$ 
12:  end if
13: end for
14: return  $\hat{G}, J$ 
```

4 Implementation

4.1 Repair Steps

Algorithm 1 gives Diagnosis + Gate. Diagnosis is the error report. Gate is the triple filter. Code builds the report and runs the checks. The model proposes field values.

The runner passes the source, document ID, allowed relations, and input triples to the model. The scorer loads the reference values. Tests change reference values and defect labels and check that model inputs stay the same. The older data loader reads the document ID from the reference record as task information.

With hash-based lookups, cleanup takes $O(|G|)$ triple operations on average. Filtering k candidates costs $O(k(1 + c_X))$ and uses $O(k + |\mathcal{R}_X|)$ extra memory. Here, c_X is the cost of finding a value in the text. Text length affects this cost. We time model calls apart from these checks.

4.2 Comparing Prompt Context and Filtering

One system prompt supports three settings: basic input, input with an error report, and input with SHACL checks. Three calls give three responses for each input. We score each response before and after filtering, giving six scores. Each before/after pair uses the same triples.

The model is `google/gemma-4-26B-A4B-it` at temperature zero. The output limit is 4,000 tokens in each setting. We report input tokens because the prompt lengths differ. Request times include retries and pacing waits. We save all responses and include invalid responses in scoring. The earlier Claude tests compare methods with their own prompts.




Fig. 3 One step of edit selection. The router gives a repair score and edit-type weights. Each model or rule edit is tried on a graph copy. Graph bounds decide which edits are allowed. The selector applies the edit with the highest positive score, saves the decision, and checks again. The test allows up to twelve steps.

4.3 SHACL Prompt Baseline

Lin et al. use five prompt parts: an introduction, detected errors, SHACL rules, graph context, and instructions (Lin et al. 2025). We use their full-rules/full-graph ($M + G$) setting for the field task. The pySHACL tool checks document IDs, allowed relations, nonempty values, and field counts. A SPARQL rule checks values against the source text.

The prompt includes the full source and returns a complete JSON graph. It groups each document’s errors into one call. We call this baseline *Lin-style SHACL context (adapted)*. RDF stores unique triples. Our shared JSON input keeps duplicate copies. We score the output before and after filtering for each prompt setting.

4.4 Choosing Edits One at a Time

Figure 3 gives the edit-selection steps in Figure 1. The profile has four scores: the share of nodes linked by edges (Q_{inc}), triple uniqueness (Q_{uniq}), rule agreement (Q_{rule}), and source support (Q_{src}). The router receives these scores $\mathbf{s}$ and four graph counts $\mathbf{g}$: nodes, triples, density, and violations. It returns repair probability p_{repair} and edit-type weights π . These weights form the scope prior. The selector uses them with quality gains, edit costs, and graph bounds to score trial edits.

A rule error or supplied candidate starts the trial checks even when $p_{\text{repair}} < \tau_{\text{repair}}$. If both are absent, the selector stops. It applies the allowed edit with the highest positive score. It keeps the graph and stops if all edit scores are zero or negative. After each edit, it computes the profile and router outputs again.

For this test, we group changes between the cleaned input and model output by relation. Each group is one edit bundle. Candidate confidence is 0.7. The run allows 12 steps. Each trial graph must meet the graph bounds. We compare learned weights with an always-on uniform prior, which gives equal weight to each edit type. Both use the default score settings.

Logs save every trial graph, profile, score, decision, and stop reason. Source matching removes whitespace from the source and value. The archive also stores the earlier logs that used different whitespace rules for the two strings. Algorithm 1 and the edit selector have separate tests.

4.5 Shared Features for Training and Use

Training and repair use the same eight input features: four quality scores, node count, triple count, density, and detected-violation count. The density uses distinct directed node pairs. Training labels describe the added defects. The model weights and feature scaler store a feature version. A version mismatch switches the router to equal weights. The document ID and allowed relations are task inputs.

The source-field check finds values after allowed field labels at the start of a document, line, or sentence. It reports a missing field when the text gives a value and the graph omits it. The check skips a field if repeated labels give conflicting values. These reports enter the violation count and candidate list. The node-link score counts connected nodes. A partial star graph can have a node-link score of 100.

Hierarchy checks use the direction of management, supervision, and membership relations. They skip relations with unclear meanings. Entity checks apply to endpoints marked as entity references. Other field values are treated as text. Each edit or bundle passes the empty-graph check. A supplied candidate receives a trial check even if the repair probability is low or the graph has no detected violation.

Let $Q \in [0, 100]$ be profile quality, h the rule-error count, $c(a)$ the edit cost, $s(a)$ its type, and $z(a)$ its confidence. The default edit score is

$$U(a) = \frac{Q(G_a) - Q(G)}{100} + 0.20 \frac{\max(0, h(G) - h(G_a))}{\max(1, h(G))} \quad (8)$$

$$- 0.35c(a) + 0.05 \log(\max(\pi_{s(a)}, 0.05) + 10^{-6}) + 0.02z(a).$$

Here, G_a is the trial graph. Delete, retype, and add costs are 0.30, 0.16, and 0.06. A bundle's cost is the sum of its edit costs. The selector applies the allowed edit with the highest positive score. The four quality scores have equal weights. Trial graphs keep all bounds already met. A score below its lower bound must stay the same or increase. A checked score can fall by at most two points, subject to these bounds. Trials also meet the density bound. Comparing pairs of triples costs $O(m^2)$ for m triples. T rounds with up to K trials per round cost $O(TKm^2)$, plus text checks. The main filter scans the candidate list once.

Table 1 Test data with added defects. The last three columns count documents.

Domain	Documents	Defects	Train	Val.	Test
Government	499	998	349	75	75
Finance	500	1,000	350	75	75
Environment	500	1,000	350	75	75
Total	1,499	2,998	1,049	225	225

A second score subtracts $0.05 \log(1/3 + 10^{-6})$ from the default score. This sets the prior term to zero for equal weights. It keeps the edit order for a fixed router but shifts the point at which scores become positive. Both scores are fixed before the saved-proposal tests.

5 Evaluation

We test repair on added defects, model-built graphs, receipt text, and contract fields. Shared prompts test the effect of context. Scores before and after filtering test the gate. Human reviewers check reference differences and actual edits. We ask three questions. **RQ1**: when does repair improve a graph? **RQ2**: what do context, filtering, and edit selection change? **RQ3**: when does an existing graph help more than fresh extraction? Table 12 gives the results.

5.1 Test Setup

Source records and added defects.

The source data contain 4,728 government records from the Shaanxi Provincial Government, 1,000 finance records from the National Financial Regulatory Administration, and 500 environment records from the Ministry of Ecology and Environment. We sample up to 500 documents from each source and build a field graph for each one.

The source text lists nonempty fields as “field name: value”. Chinese sentence marks join the fields. Each value has at most 320 characters; the service label has at most 100. The same cleaned values form the reference triples. We split documents into train, validation, and test groups at 70/15/15 with seed 42. We split by source ID before adding defects. Each original document and its changed copy belong to the same group.

Each test graph has exactly two added defects. The six types are missing triple, duplicate triple, invalid relation, reversed edge, unsupported value, and hierarchy conflict. The saved data list every clean and changed triple. The test set has 225 documents, 75 per source, and 450 defects (Table 1).

Graphs built by a model.

Claude Haiku builds a new graph for each of the same 225 test documents. It receives the source text, document ID, and allowed relations. The scorer uses the structured

fields as reference values. Ten responses have invalid JSON and are scored as empty graphs. Parse success is 95.56%. In total, 77 graphs differ from the reference, with 302 triple differences. Two reviewers check a fixed sample of 200 differences. A third reviewer resolves their disagreements (Section 5.10).

Repair methods.

All methods receive the same input graph, source text, document ID, and allowed relations.

- **Source-field Copy** reads field labels and copies values directly from the text.
- **No Repair** returns the input.
- **Rule Only** removes copies, reverses an edge that ends at the document node, and removes invalid heads and relations.
- **SHACL-style Repair** checks the document head and allowed relations using SHACL shapes (Knublauch and Kontokostas 2017).
- **Direct LLM** makes one call and returns a complete graph based on the source text.
- **ReAct-style Agent** makes one call to find errors and another to act on them (Yao et al. 2023).
- **Simple Pipeline** cleans the input, makes one call, and removes copies from the response.
- **Diagnosis + Gate** adds an error report to the prompt. It filters the output by document ID, relation, source match, duplicate count, and number of values per field.

The model-based methods use `claude-haiku-4-5-20251001` (Anthropic 2025) at temperature zero on the same endpoint, except where another model is named. All scores include failed calls. On the added-defect test, Direct LLM and Simple each have one invalid JSON response. Diagnosis + Gate and ReAct have zero.

Scores and statistical tests.

A defect is repaired when the output restores the count of each clean triple and removes its changed form. We also score correct facts kept, over-repair, triple F1, exact graph match, calls, and request time. Exact graph match checks both the triples and their copy counts.

Let G_0 be the model-built input, G_1 the repaired graph, and G^* the reference. Let d count extra and missing triples, including copies. The reduction in reference differences is

$$1 - \frac{d(G_1, G^*)}{d(G_0, G^*)},$$

averaged over the 77 inputs that differ from the reference. A negative value means the output has more differences. Fact retention and over-repair also use reference triples. Human scores measure whether reviewers find an error and accept an edit.

Confidence intervals (CIs) use 10,000 paired bootstrap samples. For matched tests, each sample selects whole documents. This keeps the two added defects of a document together. Paired sign-randomization tests use 10,000 draws and seed 42. The earlier

Table 2 Repair on 225 test documents with 450 added defects. Lower over-repair, call count, and request time are better.

Method	Defect repair	Facts kept	Over-repair	Triple F1	Exact match	Calls/doc	Time/doc
Source-field Copy	1.0000	1.0000	0.0000	1.0000	1.0000	0.00	< 0.001 s
No Repair	0.0000	1.0000	0.0000	0.7615	0.0000	0.00	< 0.001 s
Rule Only	0.3200	1.0000	0.0000	0.8472	0.0622	0.00	< 0.001 s
SHACL-style Repair	0.0000	1.0000	0.0000	0.7998	0.0000	0.00	< 0.001 s
Direct LLM	0.9244	0.9687	0.0758	0.9520	0.8267	1.00	7.25 s
ReAct-style Agent	0.8911	0.9324	0.1667	0.9213	0.5956	2.00	20.50 s
Simple Pipeline	0.9511	0.9723	0.0705	0.9634	0.8400	1.00	9.27 s
Diagnosis + Gate	0.9800	0.9932	0.0167	0.9917	0.9333	1.00	11.44 s




Fig. 4 Repair of added defects. **a**, Repair rate, triple F1, and exact match with 95% CIs from document sampling. **b**, Scores and request times for methods using one or two calls.

defect and exact-match comparisons use exact McNemar tests. The model-output difference counts also use a one-sided paired Wilcoxon test.

The first matched-context tests use unadjusted p values. Later component, receipt-context, and saved-response tests use the Holm groups stated with their tables. Each 95% CI describes one comparison. Holm correction applies to the p values within each test group.

5.2 Repair of Added Defects

Diagnosis + Gate repairs 98.00% of defects, keeps 99.32% of clean facts, and reaches 99.17% F1 (Table 2). Simple uses the same one-call limit and reaches 95.11% repair and 96.34% F1. Diagnosis + Gate repairs 15 defects that Simple misses. Simple repairs two that Diagnosis + Gate misses (McNemar $p = 0.00235$). The paired repair gain is 2.89 points (95% CI: 1.11–4.89; $p = 0.0067$). The F1 gain is 2.84 points (95% CI: 1.65–4.25). Direct LLM reaches 92.44% repair. Source-field Copy reaches 100% repair and exact match with zero model calls. It reads the field names and values directly from these records.

Diagnosis + Gate repairs 28 defects that Direct LLM misses. Direct LLM repairs three that Diagnosis + Gate misses ($p = 4.65 \times 10^{-6}$). The same counts against ReAct are 42 and two ($p = 1.13 \times 10^{-10}$). These tests compare complete repair methods. Section 5.4 tests their steps.

Table 3 Repair of model-built graphs. Difference reduction uses the 77 inputs that differ from the reference. Other scores use all 225 documents.

Method	Parse	Non-exact inputs	Difference reduction	Facts kept	Over-repair	Triple F1	Exact match
Source-field Copy	1.0000	77	1.0000	1.0000	0.0000	1.0000	1.0000
Extracted graph, no repair	0.9556	77	0.0000	1.0000	0.0000	0.8799	0.6578
Simple Pipeline	1.0000	77	-0.5915	0.9971	0.2389	0.8807	0.6711
Diagnosis + Gate	0.9956	77	0.2775	1.0000	0.0292	0.9248	0.6844




Fig. 5 Repair scores by input type. The valid-JSON group shows repair of parsed graphs. The invalid-JSON group shows recovery from empty inputs.

5.3 Repair of Model-Built Graphs

The input graphs have 87.99% F1 and 65.78% exact match (Table 3). Diagnosis + Gate raises F1 to 92.48%. It reduces reference differences by 27.75% and keeps all input triples that match the reference. Simple has a mean difference reduction of -59.15% on the 77 inputs with differences. It adds more differences than it removes.

Diagnosis + Gate leads Simple by 4.40 F1 points (95% CI: 3.07–5.77). It leaves 1.03 fewer differences per document on average (95% CI: 0.77–1.32; Wilcoxon $p = 3.18 \times 10^{-12}$). Three graphs match exactly under Diagnosis + Gate and differ under Simple. Zero show the reverse ($p = 0.25$). Source-field Copy matches all references.

Valid and invalid JSON inputs.

The 215 valid JSON inputs have 92.09% F1. Diagnosis + Gate reaches 93.53% and Simple reaches 88.85% on these inputs. Among the 67 valid inputs with reference differences, Diagnosis + Gate reduces the differences by 24.00% and raises F1 from 74.61% to 79.22%. It keeps all 148 initially exact graphs exact. The ten invalid inputs reach 69.91% F1 after repair; zero match exactly.

Table 4 Gemma prompt tests on 60 documents per input group. Each raw/gated pair uses the same response.

Context	Output	Ctrl. repair	Ctrl. F1	Natural F1	Natural exact
Base	Raw	0.9750	0.9951	0.9185	0.6667
Base	Gate	0.9750	0.9962	0.9252	0.6667
Diagnostic	Raw	0.9250	0.9850	0.9079	0.6667
Diagnostic	Gate	0.9250	0.9861	0.9147	0.6667
SHACL context	Raw	0.9500	0.9912	0.9322	0.7000
SHACL context	Gate	0.9500	0.9912	0.9364	0.7000

5.4 Prompt, Filter, and Edit-Selector Tests

We test prompt text, filtering, and edit selection in turn. On the saved Claude responses, filtering raises F1 from 98.95% to 99.17% and reduces over-repair by 0.60 points. Repair stays at 98.00% and exact match at 93.33%. The filter removes eight triples. All eight fail the source check and differ from the reference. Appendix A.1 gives the full method table.

The Gemma test uses 60 documents, 20 per source, with added-defect and model-built inputs. Three prompts share the same instructions, cleanup, temperature, and output limit. Each response is scored before and after the same gate. Table 4 gives all six settings. Table 20 gives input tokens and request times. Section 4.3 gives the SHACL prompt.

With filtering, the base prompt reaches 99.62% F1 and 97.50% repair on added-defect inputs. The error-report prompt reaches 98.61% F1 and 92.50% repair. The paired F1 change is -1.01 points (95% CI: $[-1.67, -0.38]$; $p = 0.0082$). On model-built inputs, Base + Gate reaches 92.52% F1, Diagnosis + Gate reaches 91.47%, and SHACL Context + Gate reaches 93.64%. The Diagnosis–SHACL F1 difference is -2.16 points (95% CI: $[-5.26, 0.02]$; $p = 0.105$).

On model-built inputs with an error-report prompt, filtering raises F1 from 90.79% to 91.47%. The gain is 0.69 points (95% CI: 0.20–1.29; $p = 0.0335$), with the same exact-match score. Across the 360 responses, the gate rejects 27 triples. All 27 differ from the reference.

Testing cleanup, error reports, and filtering together.

We switch input cleanup (P) and the error report (D) on or off. The four prompt settings use 60 added-defect and 60 model-built inputs, giving 480 Gemma responses. Each response is scored with and without the gate (G). Each main effect averages over the other two switches.

All six Holm-corrected p values exceed 0.05 (Table 5). The error report lowers mean F1 in both input groups. The gate raises mean F1 by 0.164 points on added-defect inputs and 0.707 points on model-built inputs. Across all input groups, the full filter rejects 47 triples: 17 have wrong document IDs and 30 fail source matching. All 47 differ from the reference. The released scores also list the joint switch effects and removals by each check.




Fig. 6 Prompt, filter, and selector tests. Panels a–b use Gemma outputs on 60 documents. Each raw/gated pair uses the same response. Bars show 95% document-bootstrap CIs. Panel c uses saved Claude proposals for 225 documents.

Table 5 Mean F1 changes in percentage points. Each input group has 60 paired documents. Each effect averages over the other two switches. CIs are 95% document-bootstrap intervals. p_H corrects the six main-effect tests with Holm correction.

Input	Component	Change	95% interval	p_H
Controlled	Preprocessing	-0.028	[-0.797, 0.778]	1.000
Controlled	Diagnosis	-0.651	[-1.348, -0.056]	0.164
Controlled	Gate	+0.164	[0.054, 0.301]	0.156
Natural	Preprocessing	+0.697	[0.000, 1.905]	1.000
Natural	Diagnosis	-0.535	[-2.376, 0.842]	1.000
Natural	Gate	+0.707	[0.225, 1.316]	0.086

Tests of each filter check.

We score 540 saved responses under eight filter settings, giving 4,320 outputs. The full filter rejects 28 triples: 11 have wrong document IDs and 17 fail source matching. All 28 differ from the reference. The other checks remove zero extra candidates in this sample. All 21 adjusted check-comparison p values exceed 0.05. Appendix A.2 gives the table.

Edit selection.

On the saved Claude proposals, the default selector with learned or equal weights accepts zero edits on added-defect inputs. Repair stays at 32.00% and F1 at 84.72%, the cleaned-input scores. In the v2 runs, model-built inputs reach 88.33% F1 with equal weights and 87.99% with learned weights. The learned trigger also reaches

Table 6 Two-model results on the same 60 added-defect inputs.

Model	Method	Repair	F1	Exact
Claude	Direct	0.958	0.970	0.867
Claude	Simple	0.958	0.966	0.833
Claude	Diag.+Gate	0.975	0.991	0.917
Gemma	Direct	0.958	0.988	0.883
Gemma	Simple	0.967	0.984	0.867
Gemma	Diag.+Gate	0.983	0.990	0.917

87.99%. Removing edit cost raises added-defect F1 to 89.30%. Model-built F1 is 88.13% in that setting. The logs show missing fields in graphs with high profile scores. Some useful additions receive scores of zero or less and are rejected. Appendix A.2 gives all 13 settings, two prior score forms, and 11,700 outcomes. Appendix A also gives earlier selector results and the call-policy test.

5.5 Tests with Two Models

We compare Direct LLM, Simple, and Diagnosis + Gate on 60 added-defect inputs, 20 per source. The models are Claude Haiku and `google/gemma-4-26B-A4B-it` (Gemma Team 2026). Each method uses the same inputs, prompt form, temperature, and one-call limit across models. Diagnosis + Gate has the highest repair rate with both models (Table 6). Its F1 lead over Simple is larger with Claude. Gemma’s direct outputs have higher F1 than Claude’s.

5.6 Source-Line Hints on New Receipts

In the first 60-receipt test, base and error-report prompts keep all input graphs unchanged. Filtering lowers F1 from 79.44% to 79.15%. It rejects an address because its spaces differ from the source. Of 49 reference differences, 27 involve case, spaces, punctuation, or amount format (Appendix A.3). We use these findings to add field meanings, source-line hints, and a shared whitespace rule.

We set the index design on 20 new receipts, then fix it before testing 60 further receipts (Section 3.4). The first 60, development 20, and new test 60 receipts have different IDs and different text after whitespace cleanup. Sampling uses seed 20260923. Test labels are downloaded after all predictions are saved. Every method receives the same field meanings, numbered source lines, input graph, and instructions. Simple uses them directly. Evidence Index adds field-line hints. SHACL adds rule-check context. All three use the same whitespace rule in the filter.

Development F1 is 80.00% for the input, 82.50% for Simple, 85.00% for Evidence Index, and 80.00% for SHACL. We keep one index design and all three methods for the test. The main comparison is Evidence Index minus Simple in mean document F1. The code and scoring rules are fixed before testing.

Table 7 Repair on 60 new receipts. Methods share field meanings, source lines, and the same whitespace check. Exact match uses the labelled fields.

Method	Triple F1	Exact graph	Normalized F1	Facts kept
Extracted input	0.7833	0.3500	0.7958	1.0000
Simple, raw	0.7875	0.2833	0.8083	0.9250
Simple, gate	0.7875	0.2833	0.8083	0.9250
Evidence Index, raw	0.8333	0.4000	0.8625	0.9306
Evidence Index, gate	0.8351	0.4000	0.8643	0.9306
SHACL context, raw	0.7833	0.2833	0.8042	0.9347
SHACL context, gate	0.7833	0.2833	0.8042	0.9347

On the new test, input F1 is 78.33%. Simple, Evidence Index, and SHACL reach 78.75%, 83.51%, and 78.33% after filtering. Simple gains 0.42 points over the input (95% CI: $[-4.17, 5.00]$). Evidence Index leads Simple by 4.76 points (95% CI: $[2.50, 7.26]$; paired randomization $p = 0.0004$). Exact match is 28.33%, 40.00%, and 28.33% for the three methods.

The input has 52 fields that differ from the reference. Simple recovers 16 of them and loses 15 exact input matches. Evidence Index recovers 26 and loses 14. SHACL recovers 13 and loses 13. The index fixes 11 punctuation differences and 6 text-boundary differences. Among its 14 lost exact matches, 10 keep the same numeric amount and change the display format.

A numeric check removes currency prefixes and thousands separators from the total field. It gives 58/60 correct amounts for the input, 59/60 for Simple, 59/60 for Evidence Index, and 58/60 for SHACL. After case and whitespace cleanup, F1 is 86.43% for Evidence Index and 80.83% for Simple. The index improves both exact and cleaned-text scores.

The filter rejects 1 candidate. It differs from the reference and joins address parts into text absent from the source. The count of rejected reference matches is 0. Filtering raises indexed F1 by 0.18 points. The other two methods keep the same F1.

Simple, Evidence Index, and SHACL change 31, 39, and 28 input graphs. Their mean repair request times are 7.91, 7.88, and 8.11 seconds. Mean input lengths are 1784, 2153, and 3404 tokens. The 240 requests produce 60 extraction and 180 repair outputs. Request times include pacing and retries. Each raw/gated score pair uses one response.

5.7 Comparing Source-Line Hints

We compare six prompt settings on the same 60 new-test receipts. The test uses 360 new Gemma responses from the same batch as the component tests. All responses parse. Each setting keeps the full source and shares the model, output limit, and filter. The random index has about the same character length as the full index.

The full index leads Simple by 4.35 F1 points (95% CI: 1.25 to 7.86). Its Holm-corrected value is $p_H = 0.065$ across five planned comparisons. The full index and anchor lines alone have the same mean F1. The full-random difference is 1.85 points (CI: $[-0.65, 4.52]$; $p_H = 0.696$). Removing field meanings from the full index lowers F1




Fig. 7 New receipt test. Left: mean F1 with 95% document-bootstrap CIs. Right: paired F1 changes. All methods share field meanings and source lines.

Table 8 Source-line tests on 60 receipts. Each row uses the same model, source, output limit, and filter. Random hints have about the same character length as the index.

Context	Exact F1	Normalized F1
Simple	0.7917	0.8125
Full index	0.8351	0.8643
anchors only	0.8351	0.8643
Random index	0.8167	0.8458
Simple, no definitions	0.8083	0.8292
Full index, no definitions	0.7470	0.7762

by 8.81 points (CI: [4.23, 13.39]; $p_H = 0.006$). The Simple setting scores higher with field meanings removed in this batch. These results show that field meanings help the indexed prompt choose values.

5.8 Repair versus Fresh Extraction

We reuse the 60 new-test receipts for four settings: repair or fresh extraction, with or without the index. All settings share source lines, field meanings, the document ID, and the output limit. Repair also receives the existing graph. All 240 responses put JSON inside Markdown fences. The fixed strict parser scores them as failures.

We then rescore the same responses after removing the fences and assigning the known document ID to every output. This extra analysis gives 83.93% F1 for indexed repair and 74.86% for indexed fresh extraction. The paired difference is 9.07 points (95% CI: 5.22 to 13.10). The repair calls use an existing graph; its creation cost is counted apart. Appendix A.4 gives the strict scores, fence-cleanup scores, ID errors, and source-line coverage. The contract test fixes fence handling and ID assignment before collection.




Fig. 8 Component and source-line tests. Points show mean paired F1 changes. Bars show 95% CIs for each change. The component and receipt tests use separate Holm groups.

Table 9 Contract results on 56 documents. F1 checks exact text after whitespace cleanup. Lost counts correct input facts removed. The input row gives the initial graph cost; repair rows give the repair-call cost.

Output	F1 (%)	Lost	New incorrect	Requests
Initial extraction	46.76	0	0	59
Simple repair	43.35	7	2	71
Indexed repair	44.84	5	12	69
Indexed re-extraction	44.12	8	32	71

5.9 Contract Field Test

We use CUAD (Hendrycks et al. 2021) to test five fields: document name, agreement date, effective date, expiration date or initial term, and governing law. We select contracts with at most one distinct reference value per field after whitespace cleanup and at most 40,000 source characters. We keep the full source and check the positions of the labelled values. Twenty official-training contracts form the development set. All 56 official-test contracts that meet the rules form the test set.

We compare initial extraction, Simple repair, indexed repair, and indexed fresh extraction. Both repair methods receive the same initial graph. The indexed methods add field-line windows to the full source. All settings share the Gemma model, field meanings, and a 4,000-token output limit. Before collection, we set the code to handle JSON fences and assign document IDs. We keep the prompts and selection rules fixed after one development round.

Indexed repair reaches 44.84% F1, compared with 46.76% for the input, a 1.92-point drop. It removes five correct input facts and adds twelve facts that differ from the reference. Simple repair also scores below the input.

Indexed repair leads indexed fresh extraction by 0.72 F1 points (95% paired CI: [-4.80, 6.48]; sign-flip $p = 0.810$). The test gives 224 outputs from 270 requests, including




Fig. 9 Contract scores. Points show mean document F1. Bars show 95% CIs for each method. The text gives the paired repair-extraction CI.

Table 10 Human scores. Agreement and κ use the two original 0/1/U answers. Positive counts use final labels. Acceptance rates use the cases with a 0 or 1 answer.

Task	Judgment	Agreement	κ	Positive / certain	U
D	Input error	80.0%	0.600	95/200 (47.5%)	0
D	Suggested edit acceptable	62.0%	0.334	86/200 (43.0%)	0
E	Input error	86.0%	0.677	134/200 (67.0%)	0
E	Actual edit acceptable	65.0%	0.128	134/199 (67.3%)	1

retries. Failed outputs are scored as empty graphs. The score checks copied text against CUAD labels after whitespace cleanup.

5.10 Human Review of Errors and Edits

Task D checks 200 reference differences sampled from the 302 triple differences in model-built inputs. Task E checks 200 actual edits sampled from 598 edits across five settings. Both use random sampling with seed 42. Each item is one difference or edit. A document can supply several items. Reviewers see the source, schema, input graph, and named operation. Task E also shows the full output graph. Method names and reference-quality hints are hidden.

Two reviewers with computer-science backgrounds and Chinese reading skills read the rules and practice examples. They then label the same 400 items on their own. For each item, they answer two questions: is the input wrong at this position (**is_error**), and is the add or remove operation acceptable (**repair_acceptable**)? They use 0, 1, or U and give reasons. U means the shown evidence is too limited for a decision. A third reviewer checks the source and graphs and resolves 154 items with a disagreement by choosing between the two answers. The 246 agreed items keep their labels. One edit-acceptance answer remains U. We retain the original labels and agreement scores.

Reviewers find input errors in 95/200 reference differences and accept 86/200 suggested edits. They accept 134/199 actual edits; one is uncertain. A missing value can

Table 11 Actual-edit acceptance by setting. The random sample has different group sizes. Rates use edits with a clear decision.

Configuration	Sampled	Acceptable	U
Archived Simple	105	63/104 (60.6%)	1
Archived Diagnosis + Gate	41	25/41 (61.0%)	0
Gemma Base + Gate	16	13/16 (81.3%)	0
Gemma Diagnosis + Gate	16	13/16 (81.3%)	0
Gemma SHACL Context + Gate	22	20/22 (90.9%)	0

be correct, yet adding it can leave two values in a field that allows one. The input-error and edit questions cover these two decisions. Original edit agreement in task E is 65.0% with $\kappa = 0.128$.

The sample has 105 Archived Simple edits, 41 Archived Diagnosis + Gate edits, 16 Gemma Base + Gate edits, 16 Gemma Diagnosis + Gate edits, and 22 Gemma SHACL Context + Gate edits. The table reports the share of accepted edits in each group.

5.11 Discussion

Source format guides the choice of repair method. A parser can copy fields from records that list names and values. Receipts require choosing values from the full text. Source-line hints raise F1 in the new receipt test. The later index tests show the role of field meanings: removing them lowers indexed F1 by 8.81 points. On contracts, indexed repair reaches 44.84% F1, compared with 46.76% for the input. Keeping correct input fields is a clear target for repair.

The tests also show what the filter changes. Source and document-ID checks remove extra triples. Field selection happens in model generation and in the first-value rule. For example, a source can contain both a subtotal and a final amount. The repair method must choose the right one for “total”. This calls for checks that use the meaning of the relation as well as the presence of the value in the text.

Human reviewers accept 67.3% of the sampled edits across five settings. They find input errors in 47.5% of the sampled reference differences. The Archived Diagnosis + Gate sample has 25/41 accepted edits. Together with the raw labels and reasons, these scores help locate errors in field choice and individual edit decisions.

5.12 Code and Data

Code, prompts, data lists, model outputs, and scoring scripts are at https://github.com/turambar928/MCP_based_KG_construction. Each run saves model settings, parse status, requests, and token use. All responses are kept. Experiment READMEs give data sources, samples, checksums, and commands.

The human-review files include fixed task rules, the scorer, summary results, original agreement scores, and sample lists. Individual answers and reasons are kept private. Receipt scores use SROIE’s human field labels.

Table 12 Research questions and test results.

Question	Comparison	Finding
RQ1	Added defects and model-built graphs	Diagnosis + Gate improves repair and F1. Direct copying matches all Chinese record references.
RQ2	Shared prompts and responses; source-line hints; edit selectors	Source and ID checks remove extra triples. The receipt index raises F1. Error reports lower mean F1 in matched tests. Equal selector weights score above learned weights on model-built inputs.
RQ3	Receipt and contract repair versus fresh extraction	The receipt test favours the existing graph after fence and ID cleanup. Indexed contract repair scores below the initial graph.

6 Conclusion

We repair document graphs with source text and field checks. The method cleans the input, asks a model for a complete graph, and filters the triples. Shared prompts and scores before and after filtering show what each step changes.

Diagnosis + Gate repairs 98.00% of the added defects and raises F1 from 87.99% to 92.48% on model-built graphs. Field copying recovers every reference in the Chinese records. Source-line hints raise receipt F1 by 4.76 points over Simple. The full index and anchors alone have the same mean F1 in the later test. Indexed contract repair reaches 44.84% F1, compared with 46.76% for the input. Its difference from fresh extraction is 0.72 points (95% CI: $[-4.80, 6.48]$). Source and document-ID checks remove unwanted triples. Error reports lower mean F1 in the matched tests, and the default learned selector scores below equal weights on model-built inputs.

Human review finds input errors in 95/200 reference differences. Reviewers accept 134/199 edits across five settings, with one uncertain edit. These results give clear targets for better repair: select the right value for each field, keep correct input facts, and check the final graph.

Statements and Declarations

Data availability

Data lists, fixed sample IDs, model outputs, and scoring scripts are at https://github.com/turambar928/MCP_based_KG_construction. Experiment READMEs give the source datasets, access links, and data changes. We share human-review totals and the fixed sampling rules. Individual answers, written reasons, and source context are kept private.

Code availability

The same repository contains the code, full prompts, model settings, and commands for running the reported analyses.

References

- Acosta M, Zaveri A, Simperl E, et al (2013) Crowdsourcing linked data quality assessment. In: International Semantic Web Conference, pp 260–276, https://doi.org/10.1007/978-3-642-41338-4_17
- Anthropic (2025) Introducing Claude Haiku 4.5. Anthropic model release announcement, URL <https://www.anthropic.com/news/claude-haiku-4-5>, published 15 October 2025; accessed 17 September 2026
- Auer S, Bizer C, Kobilarov G, et al (2007) DBpedia: A nucleus for a web of open data. In: The Semantic Web: 6th International Semantic Web Conference and 2nd Asian Semantic Web Conference (ISWC/ASWC 2007). Springer, pp 722–735, https://doi.org/10.1007/978-3-540-76298-0_52
- Bian H (2025) LLM-empowered knowledge graph construction: A survey. arXiv preprint, <https://doi.org/10.48550/arXiv.2510.20345>, [arXiv:2510.20345](https://arxiv.org/abs/2510.20345)
- Bollacker K, Evans C, Paritosh P, et al (2008) Freebase: A collaboratively created graph database for structuring human knowledge. In: Proceedings of the 2008 ACM SIGMOD International Conference on Management of Data, pp 1247–1250, <https://doi.org/10.1145/1376616.1376746>
- Bordes A, Usunier N, Garcia-Duran A, et al (2013) Translating embeddings for modeling multi-relational data. In: Advances in Neural Information Processing Systems, URL https://papers.nips.cc/paper_files/paper/2013/hash/1cecc7a77928ca8133fa24680a88d2f9-Abstract.html
- Chen M, Zhang W, Zhang W, et al (2019) Meta relational learning for few-shot link prediction in knowledge graphs. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP), pp 4217–4226, <https://doi.org/10.18653/v1/D19-1431>
- Das R, Dhuliawala S, Zaheer M, et al (2018) Go for a walk and arrive at the answer: Reasoning over paths in knowledge bases using reinforcement learning. In: International Conference on Learning Representations, URL <https://arxiv.org/abs/1711.05851>
- Dettmers T, Minervini P, Stenetorp P, et al (2018) Convolutional 2D knowledge graph embeddings. In: Proceedings of the AAAI Conference on Artificial Intelligence, <https://doi.org/10.1609/aaai.v32i1.11573>

- Ding B, Wang Q, Wang B, et al (2018) Improving knowledge graph embedding using simple constraints. In: Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics, pp 110–121, <https://doi.org/10.18653/v1/p18-1011>
- Dong X, Gabrilovich E, Heitz G, et al (2014) Knowledge vault: A web-scale approach to probabilistic knowledge fusion. In: Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, pp 601–610, <https://doi.org/10.1145/2623330.2623623>
- Galárraga LA, Teflioudi C, Hose K, et al (2013) AMIE: Association rule mining under incomplete evidence in ontological knowledge bases. In: Proceedings of the 22nd International Conference on World Wide Web, pp 413–422, <https://doi.org/10.1145/2488388.2488425>
- Gemma Team (2026) Gemma 4 technical report. arXiv preprint, <https://doi.org/10.48550/arXiv.2607.02770>, [arXiv:2607.02770](https://arxiv.org/abs/2607.02770)
- Guo S, Wang Q, Wang L, et al (2016) Jointly embedding knowledge graphs and logical rules. In: Proceedings of the 2016 Conference on Empirical Methods in Natural Language Processing, pp 192–202, <https://doi.org/10.18653/v1/d16-1019>
- Hendrycks D, Burns C, Chen A, et al (2021) CUAD: An expert-annotated nlp dataset for legal contract review. In: Advances in Neural Information Processing Systems
- Huang Z, Chen K, He J, et al (2019) ICDAR2019 competition on scanned receipt OCR and information extraction. In: 2019 International Conference on Document Analysis and Recognition (ICDAR). IEEE, pp 1516–1520, <https://doi.org/10.1109/ICDAR.2019.00244>
- Ji S, Pan S, Cambria E, et al (2022) A survey on knowledge graphs: Representation, acquisition, and applications. IEEE Transactions on Neural Networks and Learning Systems 33(2):494–514. <https://doi.org/10.1109/tnnls.2021.3070843>
- Knublauch H, Kontokostas D (2017) Shapes constraint language (SHACL). W3C Recommendation, URL <https://www.w3.org/TR/2017/REC-shacl-20170720/>
- Kontokostas D, Westphal P, Auer S, et al (2014) Test-driven evaluation of linked data quality. In: Proceedings of the 23rd International Conference on World Wide Web, pp 747–758, <https://doi.org/10.1145/2566486.2568002>
- Lao N, Mitchell T, Cohen WW (2011) Random walk inference and learning in a large scale knowledge base. In: Proceedings of the 2011 Conference on Empirical Methods in Natural Language Processing, pp 529–539, URL <https://aclanthology.org/D11-1049/>

- Lewis P, Perez E, Piktus A, et al (2020) Retrieval-augmented generation for knowledge-intensive NLP tasks. In: Advances in Neural Information Processing Systems, pp 9459–9474, URL https://papers.nips.cc/paper_files/paper/2020/hash/6b493230205f780e1bc26945df7481e5-Abstract.html
- Lin TW, Fierro G, Li H, et al (2025) Systematic evaluation of knowledge graph repair with large language models. arXiv preprint, <https://doi.org/10.48550/arXiv.2507.22419>, arXiv:2507.22419
- Neo4j Labs (2024) LLM Knowledge Graph Builder. Open-source software repository, URL <https://github.com/neo4j-labs/llm-graph-builder>, version inspected at commit 5ff7af3e9bb9226e1bbeed02f70f8d98697727a7; accessed 16 September 2026
- Nickel M, Murphy K, Tresp V, et al (2016) A review of relational machine learning for knowledge graphs. Proceedings of the IEEE 104(1):11–33. <https://doi.org/10.1109/jproc.2015.2483592>
- Ortona S, Meduri VV, Papotti P (2018) Robust discovery of positive and negative rules in knowledge bases. In: Proceedings of the 34th IEEE International Conference on Data Engineering, pp 1168–1179, <https://doi.org/10.1109/icde.2018.00108>
- Pan JZ, Razniewski S, Kalo JC, et al (2023) Large language models and knowledge graphs: Opportunities and challenges. arXiv preprint, <https://doi.org/10.48550/arXiv.2308.06374>, arXiv:2308.06374
- Pan S, Luo L, Wang Y, et al (2024) Unifying large language models and knowledge graphs: A roadmap. IEEE Transactions on Knowledge and Data Engineering 36(7):3580–3599. <https://doi.org/10.1109/tkde.2024.3352100>
- Paulheim H (2017) Knowledge graph refinement: A survey of approaches and evaluation methods. Semantic Web 8(3):489–508. <https://doi.org/10.3233/sw-160218>
- Paulheim H, Bizer C (2014) Improving the quality of linked data using statistical distributions. International Journal on Semantic Web and Information Systems 10(2):63–86. <https://doi.org/10.4018/ijswis.2014040104>
- Qu M, Tang J (2019) Probabilistic logic neural networks for reasoning. In: Advances in Neural Information Processing Systems, URL https://papers.nips.cc/paper_files/paper/2019/hash/13e5ebb0fa112fe1b31a1067962d74a7-Abstract.html
- Qu M, Chen J, Xhonneux LP, et al (2021) RNNLogic: Learning logic rules for reasoning on knowledge graphs. In: International Conference on Learning Representations, URL <https://arxiv.org/abs/2010.04029>
- Schlichtkrull M, Kipf TN, Bloem P, et al (2018) Modeling relational data with graph convolutional networks. In: European Semantic Web Conference, pp 593–607, https://doi.org/10.1007/978-3-319-93417-4_38

- Shi B, Weninger T (2017) ProjE: Embedding projection for knowledge graph completion. In: Proceedings of the AAAI Conference on Artificial Intelligence, <https://doi.org/10.1609/aaai.v31i1.10677>
- Suchanek FM, Kasneci G, Weikum G (2007) YAGO: A core of semantic knowledge. In: Proceedings of the 16th International Conference on World Wide Web, pp 697–706, <https://doi.org/10.1145/1242572.1242667>
- Wang Q, Mao Z, Wang B, et al (2017) Knowledge graph embedding: A survey of approaches and applications. *IEEE Transactions on Knowledge and Data Engineering* 29(12):2724–2743. <https://doi.org/10.1109/tkde.2017.2754499>
- Wang X, Chen L, Ban T, et al (2021) Knowledge graph quality control: A survey. *Fundamental Research* 1(5):607–626. <https://doi.org/10.1016/j.fmre.2021.09.003>
- Wei J, Wang X, Schuurmans D, et al (2022) Chain-of-thought prompting elicits reasoning in large language models. In: *Advances in Neural Information Processing Systems*, pp 24824–24837, <https://doi.org/10.52202/068431-1800>
- White J, Fu Q, Hays S, et al (2023) A prompt pattern catalog to enhance prompt engineering with ChatGPT. In: *Proceedings of the 30th Conference on Pattern Languages of Programs*. ACM, <https://doi.org/10.64346/PLoP2023p05>
- Wienand D, Paulheim H (2014) Detecting incorrect numerical data in DBpedia. In: *European Semantic Web Conference*, pp 504–518, https://doi.org/10.1007/978-3-319-07443-6_34
- Yang B, Yih Wt, He X, et al (2015) Embedding entities and relations for learning and inference in knowledge bases. In: *International Conference on Learning Representations*, URL <https://arxiv.org/abs/1412.6575>
- Yao S, Zhao J, Yu D, et al (2023) ReAct: Synergizing reasoning and acting in language models. In: *International Conference on Learning Representations*, URL <https://arxiv.org/abs/2210.03629>
- Zaveri A, Rula A, Maurino A, et al (2016) Quality assessment for linked data: A survey. *Semantic Web* 7(1):63–93. <https://doi.org/10.3233/sw-150175>
- Zhu Z, Zhang Z, Xhonneux LP, et al (2021) Neural Bellman-Ford networks: A general graph neural network framework for link prediction. In: *Advances in Neural Information Processing Systems*, pp 29476–29490, URL https://papers.nips.cc/paper_files/paper/2021/hash/f6a673f09493afcd8b129a0bcf1cd5bc-Abstract.html

Table 13 Four settings on added-defect inputs. Direct LLM uses a different prompt.

Variant	Repair	Over-repair	F1	Exact
No LLM completion	0.3200	0.0000	0.8472	0.0622
Direct LLM (separate prompt)	0.9244	0.0758	0.9520	0.8267
No constraint gate	0.9800	0.0227	0.9895	0.9333
Diagnosis + Gate	0.9800	0.0167	0.9917	0.9333

A Extra Tests

This section gives earlier method results, full filter and selector runs, the first receipt test, parsing details, and extra checks. Section 5 gives the main results.

A.1 Earlier Method and Filter Tests

Table 13 compares four saved settings. The no-completion setting returns the cleaned graph. Direct LLM uses its own prompt and skips cleanup. The filter test scores the same proposal before and after checks. Section 5.4 tests prompt text. Section A.2 tests each filter and selector step.

Model completion recovers most missing fields. Diagnosis + Gate reaches 99.17% F1, compared with 95.20% for Direct LLM. Filtering keeps the same repair rate and reduces over-repair by 0.60 points. It removes eight triples that fail source matching. All eight differ from the reference.

A.2 Tests of Filter and Selector Steps

We use 360 Gemma responses from the shared-prompt test and 180 responses from the new receipt test. We score each under eight filter settings: all checks, each of five checks removed, duplicate and field-count checks removed together, and all checks removed. The 540 responses give 4,320 scored outputs. Candidate order stays fixed. Removing two checks together tests how they affect each other.

Table 14 gives full-filter F1 minus each changed-filter F1. We first average changes across the three prompts for each document, then across documents. The full filter rejects 28 triples. Eleven have a wrong document ID and 17 fail source matching. All 28 differ from the reference. Other checks remove zero extra candidates. On model-built inputs, source matching adds 0.505 F1 points and the ID check adds 0.085 points. The 21 adjusted test p values exceed 0.05. The logs also list correct facts lost when a check is removed.

Current edit-selector tests.

We run the v2 selector on saved proposals for 225 added-defect and 225 model-built inputs. Thirteen settings give 5,850 outcomes with the default absolute prior score. We repeat them with a score relative to equal weights, giving another 5,850 outcomes. Both runs use `public-profile-v2` features and the retrained router. The archive keeps the older v1 runs.

Table 14 F1 change from each filter check, in percentage points. Each input group has 60 documents and three prompts. Positive values favour the full filter.

Removed check(s)	Controlled	Natural	Receipts
Duplicate	+0.000	+0.000	+0.000
Document head	+0.037	+0.085	+0.000
Relation type	+0.000	+0.000	+0.000
Source support	+0.037	+0.505	+0.060
Field count	+0.000	+0.000	+0.000
Duplicate + field count	+0.000	+0.000	+0.000
All checks	+0.073	+0.590	+0.060

Table 15 V2 selector with the default absolute prior score. Accepted counts are actions or bundles, shown as added-defect / model-built. Each F1 uses 225 documents.

Setting	Controlled F1	Natural F1	Accepted
Always-on uniform	0.8472	0.8833	0 / 10
Learned prior, always on	0.8472	0.8799	0 / 10
Learned prior and trigger	0.8472	0.8799	0 / 10
No prior penalty	0.8472	0.8833	0 / 10
No action cost	0.8930	0.8813	68 / 13
No density bound	0.8472	0.8833	0 / 10
No quality bounds	0.8472	0.8833	0 / 10
One iteration	0.8472	0.8833	0 / 10
No rule candidates	0.8472	0.8833	0 / 9
Rule candidates only	0.8472	0.8799	0 / 10
No local score/bounds	0.8472	0.8799	0 / 0
No graph score/bound	0.8472	0.8833	0 / 10
No source score/bound	0.8472	0.8799	0 / 10

The base setting uses equal weights and keeps the selector on. Each test changes one named step. The trigger test starts from the always-on learned-weight setting. Removing a score sets its term to zero and removes its bounds. Other weights stay the same. Rule detectors and violation rewards stay active. Removing rule candidates drops detector actions and direct source-field additions. The graph profile and saved model proposals stay in use.

The default equal-weight selector reaches 84.72% F1 on added-defect inputs and accepts zero edits. Removing edit cost raises F1 to 89.30% with 68 accepted actions. On model-built inputs, equal weights reach 88.33% F1. Learned weights and the learned trigger reach 87.99%. Every setting keeps all initially correct reference facts in both groups. The released files give document scores, trial edits, rejection reasons, and stop reasons. Holm correction groups the 24 component tests for each prior score form.

Table 16 V2 selector on saved proposals. Relative scores use equal weights as the zero point. Accepted counts are added-defect / model-built.

Prior / offset	Controlled F1	Natural F1	Accepted
Uniform / absolute	0.8472	0.8833	0 / 10
Uniform / relative	0.8472	0.8833	0 / 10
Learned / absolute	0.8472	0.8799	0 / 10
Learned / relative	0.8929	0.8764	156 / 51

Router training and prior scores.

The v2 router has an $8 \rightarrow 32 \rightarrow 16$ network, one repair output, and three edit-type outputs (884 parameters). The feature scaler uses 2,098 training rows. Validation and test each have 450 rows grouped by document. Repair and edit-type classification on added-defect inputs both reach 100%. Training labels follow the known added defects.

Table 16 gives four weight/score settings from the earlier v2 test. The full runs above include these settings. With learned weights, the relative score raises added-defect F1 from 84.72% to 89.29%. Model-built F1 falls from 87.99% to 87.64%. Both equal-weight settings return the same graphs. The absolute score remains the default.

A.3 First Receipt Test

We use SROIE receipt text and human field labels (Huang et al. 2019). A random sample selects 60 of 626 documents from the corrected trainval mirror (seed 20260922). We join text lines in their stored order. The model receives this text and four field names: company, date, address, and total. The scorer reads the field labels. One receipt has no address label, leaving 239 scored fields. Every method is scored on these fields. The model receives all four field names.

Gemma first extracts a graph. Three repair prompts receive this same graph with the settings in Section 4.2. We fix samples and settings before calling the model. The test has 60 extraction calls and 180 repair outputs. Each repair response is scored before and after filtering. We also run source copying and rule repair with their original settings. The main score uses exact triples. A second score ignores case and whitespace in values.

The input reaches 79.44% F1. After filtering, base and error-report prompts both reach 79.15%. SHACL context reaches 79.56%. The error-report minus base difference is 0.00 points (95% CI: [0.00, 0.00]; paired randomization $p = 1.0000$). Field copying reaches 0.00% F1 on receipts and 100% on the named-field Chinese records.

Of 239 reference values, 220 occur exactly in the source text. The other 19 differ from the source strings. Filtering rejects 3 candidates, all matching reference values. They use the same address in three settings. The source has a double space; the parser replaces it with a single space in the candidate. After the run, we apply the same whitespace rule to the source. All three settings then keep the address. F1 is 79.44% for base and error-report prompts and 79.86% for SHACL. Exact graph match stays at 33.33%. Table 17 gives the original-rule scores.

Table 17 First SROIE test on 60 receipts. Exact match uses labelled fields. Normalized F1 ignores case and whitespace in values.

Method	Triple F1	Exact match	Normalized F1	Facts kept
Extracted input	0.7944	0.3333	0.8153	1.0000
Rule Only	0.7944	0.3333	0.8153	1.0000
Source-field Copy	0.0000	0.0000	0.0000	0.0000
Base, raw	0.7944	0.3333	0.8153	1.0000
Base, gate	0.7915	0.3333	0.8123	0.9944
Diagnostic, raw	0.7944	0.3333	0.8153	1.0000
Diagnostic, gate	0.7915	0.3333	0.8123	0.9944
SHACL context, raw	0.7986	0.3333	0.8194	1.0000
SHACL context, gate	0.7956	0.3333	0.8165	0.9944




Fig. 10 First receipt scores. Points show mean F1 and 95% document-bootstrap CIs. Each raw/gated pair uses one response. The dashed line shows input F1.

The run has 240 recorded requests, including retries, and 240 parsed outputs. Mean repair request times are 7.87, 7.92, and 8.05 seconds for base, error-report, and SHACL prompts. Mean input lengths are 661, 715, and 2276 tokens. Request times include pacing. Extraction and context setup are counted apart.

A.4 Receipt Repair and Fresh-Extraction Details

We collect 240 further Gemma responses on the same sixty new-test receipts. The settings are repair or fresh extraction, with or without the index. All share source lines, field meanings, document ID, and output limit. Repair also receives the existing graph. The main score difference is indexed repair minus indexed fresh extraction after filtering. All 240 responses have JSON inside Markdown fences. The fixed strict parser marks them as failures and scores every setting at zero F1.

Table 18 Extra scores after collection on sixty receipts. Fence-compatible removes outer Markdown fences. ID-normalized also assigns the known document ID before filtering. Both use the saved responses.

Method	Fence-compatible F1	ID-normalized F1	Wrong-ID graphs
Fresh extraction	0.6333	0.7256	18/60
Indexed extraction	0.6105	0.7486	28/60
Repair	0.7917	0.7917	0/60
Indexed repair	0.8393	0.8393	0/60

After collection, we remove each outer fence and apply the same schema checks. All 240 responses then parse. Table 18 gives these extra scores. Indexed repair reaches 83.93% F1 and indexed fresh extraction reaches 61.05%. The paired difference is 22.88 points (95% CI: 14.84 to 31.35).

Indexed fresh extraction gives the wrong document ID in 28 graphs. Indexed repair uses the right ID in all sixty. Assigning the known ID to all four settings raises indexed fresh-extraction F1 to 74.86%. Indexed repair leads by 9.07 points (95% CI: 5.22 to 13.10). This ID change keeps field values fixed. The original 22.88-point gap includes both ID and field-value effects. Initial graph creation is counted apart from these repair calls.

Fields found by the line index.

We locate the original follow-up’s 240 reference fields in the source lines. Matching keeps case and punctuation and cleans whitespace. For a value spread over several lines, the index must include all lines of one occurrence. Of the 240 values, 188 occur in index lines or nearby lines, 36 occur outside them, and 16 differ from the source text. Among the 188 covered values, Simple returns 163 exactly and Evidence Index returns 174. Both return 26 of the 36 outside values. They return zero of the sixteen values that differ from the source. The index’s net gain of eleven exact fields comes from the covered group.

A.5 Calling Repair after Visible Errors

We test a call rule based on visible ID, relation, duplicate, field-count, and source-match errors. It selects either the input or the saved repair output for each document. On the 450-item model-built/clean stream, it calls repair on 6.22% of inputs. It misses 63.64% of graphs with reference differences. Its F1 is 94.46%, compared with 95.82% for always calling repair. The visible checks miss many missing or wrongly assigned field values. The main method calls repair on every document.

Useful edits in the earlier selector.

We use saved proposals and trial graphs from the earlier selector. These runs came before the shared-feature and candidate-start changes in Table 16. For each document, we try every subset of model edit bundles and score its graph against the reference. We apply the saved edits with stopping, cost, and graph-bound checks turned off. We

Table 19 Earlier selector results on 225 inputs per group. The oracle uses references to find the best saved subset. One document can have several trials.

Diagnostic	Controlled	Natural
Preprocessed graph F1	0.8472	0.8799
Uniform optimizer F1	0.8472	0.8833
Fixed-proposal oracle F1	0.9964	0.9173
Initially no detected violation	86	198
Imperfect among those graphs	72	50
Improvement available among those graphs	70	8
Improving trials with non-positive utility	212	1

call the best reference-scored subset the oracle. This analysis uses model bundles and reference values.

We try 2,144 subsets across 450 inputs. Table 19 lists the results. On added-defect inputs, 86 graphs have zero detected violations. Of these, 72 miss reference fields and 70 have a proposal subset that raises F1. Among the checked trials, 212 raise reference F1 but have scores of zero or less. The matching model-built counts are 8 graphs and 1 trial.

Removing edit cost raises F1 on 67 added-defect documents, keeps it the same on 158, and lowers it on 0. On model-built inputs, it raises F1 on 0, keeps it the same on 222, and lowers it on 3. The three graphs with lower F1 already differ from the reference after cleanup. All 148 exact model-built inputs keep their F1. The paired added-defect gain is 4.58 points (95% document-bootstrap CI: 3.69–5.54). The same cost change has different effects on the two input groups.

A.6 Time and Memory for Graph Checks

We time the report in Eq. (2) on separate document batches with 1K–50K triples. A full pass checks every document. A local update checks one changed document. Full times are medians of seven runs. Local times are medians of 200 runs. Table 21 gives the values. Memory measures the peak new Python memory for building the batch. Source text and schema objects are shared. Model-call times are recorded apart.

A.7 Extra Quality and Extraction Checks

A Gemma judge scores 180 source–triple pairs with method names hidden. Its source score has Pearson $r = 0.794$ and Spearman $\rho = 0.807$ with exact reference validity (both $p < 2.4 \times 10^{-40}$). A fixed 50-item sample has the same mean, 0.96, in five further runs.

We compare our extractor with Neo4j Labs’ LLMGraphTransformer core on 45 balanced TNEWS documents (Neo4j Labs 2024). Category accuracy is 0.556 for our field-based extractor and 0.533 for Graph Builder. The paired test gives $p = 1.00$.

A.8 Request Costs and Check Times




Fig. 11 Time and memory for reports on separate document batches. Local updates check one changed document. Memory measures new batch storage with shared source text.

Table 20 Request costs. Mean times include retries and pacing. Source setup and queue time are counted apart.

Input	Context	Calls	Input tokens	Output tokens	Wall (s)
Controlled	Base	1.00	947	517	8.58
Controlled	Diagnostic	1.00	1017	512	8.61
Controlled	SHACL	1.00	3080	513	8.45
Natural	Base	1.00	832	458	7.51
Natural	Diagnostic	1.00	901	462	7.90
Natural	SHACL	1.00	2822	455	8.16

Table 21 Time and memory for graph reports. Memory counts peak new batch storage in MiB. Source texts are shared.

Triples	Documents	Full (ms)	Local (ms)	Memory (MiB)
1,000	143	1.20	0.009	0.29
5,000	716	6.08	0.008	1.47
10,000	1,434	13.90	0.008	2.96
50,000	7,173	88.10	0.008	14.84